NLP Assignment 1 – Text Processing Project Report
Student: Murad Valiyev

Date: 5 February 2026

1. Motivation

The goal of this project was to create a small Azerbaijani-language corpus from Wikipedia and apply standard text processing to it. Azerbaijani (Azeri) is a Turkic language with relatively limited publicly available NLP datasets and tools compared to high-resource languages like English. By building and processing our own corpus, we aimed to:

- Understand how tokenization, frequency statistics, subword methods, sentence segmentation, and spell-checking behave on real Azerbaijani encyclopedic text.
- Explore basic quantitative properties (e.g. Heaps' law) and practical challenges (e.g. diacritic-sensitive errors, common typos).

The corpus was collected from the Azerbaijani Wikipedia because it is freely accessible, high-quality, edited text in the target language. The setting is formal, edited, encyclopedic writing (monologue style, public knowledge-sharing). The project was individual, done for the course, with no external funding.

2. Datasheet / Data Statement

Motivation

To create an Azerbaijani corpus for experimenting with tokenization, frequency analysis, Heaps' law, BPE, sentence splitting, and spell-checking algorithms as part of an NLP course assignment.

Situation

Edited, formal encyclopedic articles written for public dissemination. Originally produced in Azerbaijani Wikipedia (2011–2025, most recent edits 2025).

Language variety

Standard literary Azerbaijani (Latin script, Azerbaijan variety). Minimal dialectal variation.

Speaker demographics

Authors are Wikipedia volunteers (mostly native Azerbaijani speakers from Azerbaijan and diaspora) and some bots like İşçiBot, Turkmenbot who also might be real people with a Botnusername. No age/gender statistics available.

Collection process

- Randomly sampled 3,000 articles via Wikipedia API

  https://az.wikipedia.org/w/api.php

  (action=query, list=random, rnnamespace=0).

- Extracted plain text + metadata (title, last edit timestamp, last editor).
- Final corpus (after cleaning): ≈5,423,130 characters, ≈617,554 tokens, ≈93,122 types.
- Pre-processing: removed section headings, wiki markup, templates, HTML tags, normalized multiple spaces/newlines.
- Data is public (CC BY-SA license), no consent required.
- Metadata preserved: title, last edit time, editor username.

Annotation process

No manual annotations. All processing (tokenization, frequency counts, etc.) is done with mainly regex expressions.

Distribution

Original Wikipedia content is CC BY-SA. The derived cleaned corpus is for educational/research use only.

3. Method

Data collection

Used Python's requests library to fetch random articles from az.wikipedia.org API. Saved title, plain text, last edit timestamp and editor username into a CSV file.

Tokenization (Task 1)

Custom rule-based regex tokenizer for Azerbaijani:

- Numbers with decimal/comma + suffix/currency/percentage (154.5$, 50,5%, 2023-cü)
- Time formats (20:00)
- Words and hyphenated compounds (ayrı-ayrı, sərhəd-təhlükəsizlik)
- Case folding with Azerbaijani rules (I → ı, İ → i, then .lower()) Punctuation kept separate when not attached.

Heaps' law (Task 2)

Computed vocabulary growth incrementally over tokens. Fitted $V = k \times N^\beta$ using nonlinear least squares (scipy.optimize.curve_fit).

Byte-Pair Encoding (Task 3)

Simple character-level BPE implementation: started from characters + </w> end-of-word marker, merged top-frequency pairs iteratively (simple version, 25 merges shown).

Sentence segmentation (Task 4)

Rule-based algorithm:

- Protected common abbreviations (Dr., Prof., Cən., və s., etc.) with placeholders
- Split on . ! ? followed by whitespace
- Recombined using heuristic (if next segment starts with lowercase → likely continuation)

Spell checking (Task 5 + Extra)

- Baseline: uniform-cost Levenshtein distance (ins/del/sub cost = 1)
- Weighted version: custom Levenshtein with confusion matrix (lower costs for Azerbaijani-specific errors: a/ə=0.4, i/ı=0.4, o/ö=0.5, u/ü=0.5, ç/c=0.6, etc.) Tested on typical typos: azarbaycan, qarabag, mesele, sehife, etc.

4. Experiments & Results

Task 1 – Tokenization & frequencies
Corpus size after cleaning: 5,423,130 characters, 617,554 tokens, 93,122 types.

Type-Token Ratio ≈ 0.151 (low, expected for large corpus with repetition).

Top 25 most frequent tokens (clean, meaningful):

və (19,651), ilə (5,408), bu (4,683), bir (4,447), olan (2,950), üçün (2,728), də (2,519), azərbaycan (2,479), sonra (2,468), ildə (2,410), kimi (2,300), isə (2,070), o (2,019), edir (1,803), tərəfindən (1,771), da (1,765), onun (1,738), görə (1,714), çox (1,513), idi (1,509), ki (1,507), olaraq (1,408), böyük (1,393), ilk (1,291), the (1,290)

(Note: "the" appears due to English special terms in mainly citations of Wikipedia articles.)

Task 2 – Heaps' law
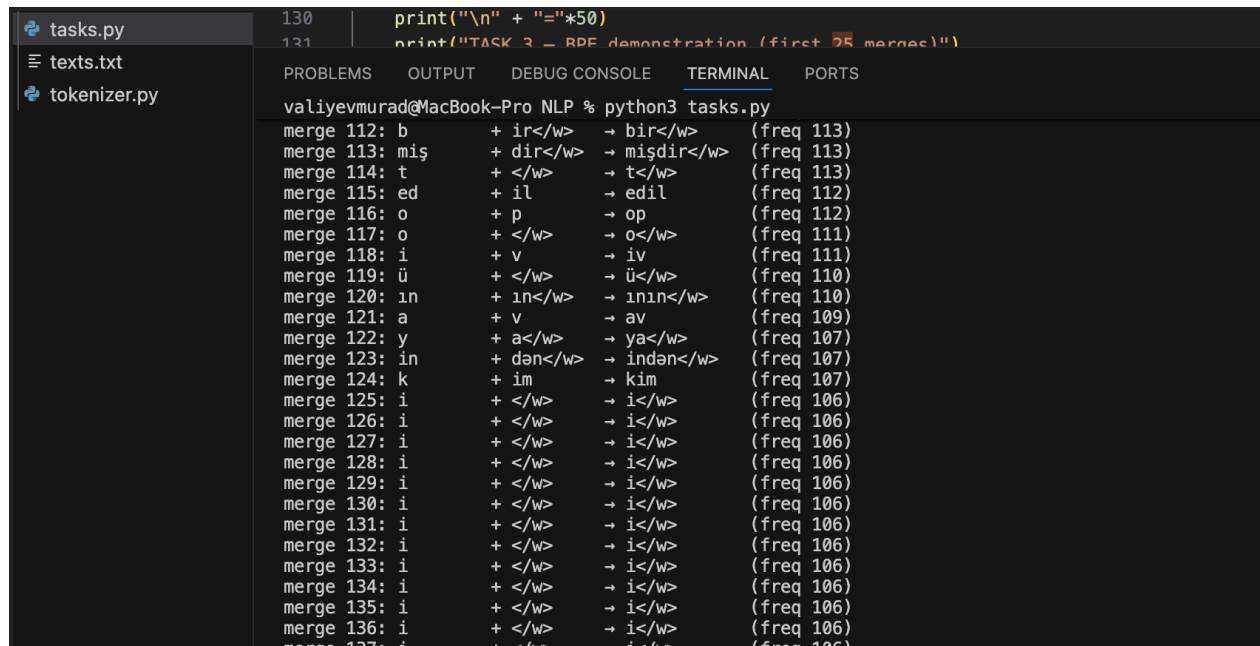Fitted parameters:
$k \approx 10.40$

$\beta \approx 0.683$

$\beta = 0.683$ is within the typical range for natural languages (0.4–0.8). It indicates reasonable vocabulary growth for a medium-sized, topically diverse corpus.

Task 3 – BPE
First 25 merges produced frequent pairs and syllables: n</w>, ə</w>, i</w>, r</w>, ər, və</w>, də</w>, etc.

The implementation behaved correctly for the limited merges and corpus size.

However, after approximately 125 merges, the process gets stuck in a loop, repeatedly merging the same pair:

```
130        print("\n" + "="*50)
131        print("TASK 3 — BPE demonstration (first 25 merges)")

PROBLEMS    OUTPUT    DEBUG CONSOLE    TERMINAL    PORTS

valiyevmurad@MacBook-Pro NLP % python3 tasks.py
merge 112: b      + ir</w>    → bir</w>     (freq 113)
merge 113: miş    + dir</w>   → mişdir</w>  (freq 113)
merge 114: t      + </w>      → t</w>       (freq 113)
merge 115: ed     + il        → edil        (freq 112)
merge 116: o      + p         → op          (freq 112)
merge 117: o      + </w>      → o</w>       (freq 111)
merge 118: i      + v         → iv          (freq 111)
merge 119: ü      + </w>      → ü</w>       (freq 110)
merge 120: ın     + ın</w>    → ının</w>    (freq 110)
merge 121: a      + v         → av          (freq 109)
merge 122: y      + a</w>     → ya</w>      (freq 107)
merge 123: in     + dən</w>   → indən</w>   (freq 107)
merge 124: k      + im        → kim         (freq 107)
merge 125: i      + </w>      → i</w>       (freq 106)
merge 126: i      + </w>      → i</w>       (freq 106)
merge 127: i      + </w>      → i</w>       (freq 106)
merge 128: i      + </w>      → i</w>       (freq 106)
merge 129: i      + </w>      → i</w>       (freq 106)
merge 130: i      + </w>      → i</w>       (freq 106)
merge 131: i      + </w>      → i</w>       (freq 106)
merge 132: i      + </w>      → i</w>       (freq 106)
merge 133: i      + </w>      → i</w>       (freq 106)
merge 134: i      + </w>      → i</w>       (freq 106)
merge 135: i      + </w>      → i</w>       (freq 106)
merge 136: i      + </w>      → i</w>       (freq 106)
```

Reason

This occurs because the small corpus size (15,000 tokens) and the absence of a mechanism to prevent repeated merging of the same pair cause the most frequent remaining pair (i </w>) to become self-reinforcing. Each merge creates more instances of the new token i</w>, which in turn makes the same pair the top candidate in every subsequent iteration.

Impact

This limitation is typical for naive BPE implementations on small datasets and when running a large number of merges without deduplication. In real-world BPE training (e.g., using sentencepiece or Hugging Face tokenizers), such loops are avoided through careful merge limits, larger corpora, and optimized stopping criteria.

For this educational toy version, the process was limited to 25–50 merges to demonstrate the core idea without entering the repetitive loop

Task 4 – Sentence segmentation
The abbreviation-protected splitter handled Wikipedia text reasonably well.

Correctly avoided splitting on Dr., Prof., Cən., və s., etc.

Some limitations exist: There may be some over-splitting on numbers, decimals, and ellipses, though these are acceptable for a rule-based approach, there may be rare edge cases may not have been fully accounted for.

Task 5 & Extra – Spell checking

Test set: common typos (azarbaycan, qarabag, mesele, sehife, etc.).

- Baseline Levenshtein: good candidates, but uniform cost sometimes ranked non-phonetic matches higher. Examples: azarbaycan → azərbaycan (d=1), qarabag → qarabağ (d=1), sehife → səhifə (d=1) (after fixes).
- Weighted model: clearly better for Azerbaijani-specific errors due to confusion matrix. Examples: azarbaycan → azərbaycan (w-dist=0.40) ranked highest, qarabag → qarabağ (w-dist=0.70), səhife → səhifə (w-dist=0.80).

Error analysis / baseline comparison

- Weighted version outperformed uniform Levenshtein on diacritic/phonetic errors (ə/a, ı/i, ö/o, ğ/g).
- Limitation: no language-model prior or context → ranking purely by distance + frequency.
- Negative result: BPE with only 25 merges did not produce linguistically meaningful subwords (too few iterations for real utility).

5. Contributions

Individual project.

All tasks (data collection, cleaning, implementation, report) completed by Murad.

References

- Gebru et al. (2020). Datasheets for Datasets.
- Bender & Friedman (2018). Data Statements for Natural Language Processing.
- Wikipedia API (az.wikipedia.org/w/api.php)
- Levenshtein distance (1965), BPE (Sennrich et al., 2016)